

# Khushi Malik

Email: [khushimalik511263@gmail.com](mailto:khushimalik511263@gmail.com) | +91- 9536110472

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

## EDUCATION

Bachelor of Science in Mathematics, Kirori Mal College, University of Delhi  
Current CGPA: 8.27/10.0

August 2023 – May 2027  
Delhi, India

## CORE COMPETENCIES

- **Programming & Data:** Python, Machine Learning, SQL Deep Learning, Natural Language Processing (NLP), Computer Vision, Data Analysis, Exploratory Data Analysis (EDA), Matplotlib
- **Tools & Frameworks:** Jupyter Notebook, Git, GitLab, PyTorch, AWS, Docker, FastAPI
- **Focus Areas:** Data Science
- **Specialized Skills:** Retrieval Augmented Generation (RAG), Model Fine-tuning, Prompt Engineering

## PROFESSIONAL EXPERIENCE

### Lead AI Researcher — Planto AI

Remote, New Delhi , August 2025 – Present

- ➔ **Streamlined Retrieval-Augmented Generation (RAG) pipelines**, elevating Large Language Model accuracy by **18%** and reducing response latency by **25%** through advanced tokenization and prompt compression techniques.
- ➔ **Led a cross-functional team of 5 engineers** in building an AI-powered VS Code extension, increasing developer productivity by **30%** through natural language coding assistance.
- ➔ **Fine-tuned Large Language Models** using HuggingFace Transformers and PEFT methods, **cutting** computational costs by 40% while maintaining model performance.
- ➔ **Presented technical findings** to stakeholders, contributing to strategic decisions that accelerated AI project adoption by **15%**.

### AI Researcher — Planto

Remote, New Delhi , May 2025 – July 2025

- ➔ **Researched and applied tokenization optimization and prompt compression** for Large Language Models, improving model inference speed by **20%** and increasing response relevance by **15%**.
- ➔ **Engineered** end-to-end ML pipelines using Python and HuggingFace, enabling **production deployment within 4 weeks** and supporting **real-time data processing for 10K+ queries per day**.
- ➔ **Collaborated** to **deploy** RAG systems, **enhancing** retrieval accuracy by 18%

### Freelance Data Scientist — Zindi

Remote, South Africa , April 2025 – Present

- ➔ Demonstrating expertise in **Computer Vision and Machine Learning** applied to large-scale agricultural datasets.
- ➔ **Processed and engineered** features on 10,000+ samples, **boosting** model accuracy by 15% and **reducing** training time by 20%.
- ➔ **Refined** ML models for agriculture and healthcare, **attaining** up to 92% predictive accuracy.

## PROJECTS

### Cogni-Compliance: End-to-End Agentic RAG System | [GitHub](#) | [Live Demo](#) | [Documentation](#)

September 2025

AI Chat Application for Legal & Regulatory Documents

- ➔ Built an agentic RAG pipeline with FAISS + BM25 hybrid retrieval and cross-encoder re-ranking for accurate Q&A over GDPR, CCPA, and HIPAA.
- ➔ Designed an agentic workflow with LangGraph, integrating Llama 3.1 for grading and Mixtral for answer synthesis.
- ➔ Developed and deployed a full-stack MERN + FastAPI system, with RAG backend on AWS EC2 and frontend/backend on Render.
- ➔ Processed legal PDFs (extraction, cleaning, semantic chunking, vectorization) using Colab/Jupyter to create a FAISS knowledge base.

### Codex-Finetune: Intelligent Coding Copilot | [Link](#)

June 2025

- ➔ Fine-tuned starcoder 2-3b model on CodeSearchNet dataset using Low-Rank Adaptation (LoRA) techniques
- ➔ Implemented multi-task code generation for functions, docstrings, and code completion using HuggingFace Transformers
- ➔ Secured final training loss of 1.20 across 3,000 training steps with Weights and Biases monitoring
- ➔ Tech Stack: Python, HuggingFace, Parameter Efficient Fine-Tuning, Machine Learning

### Amini GeoFM Decoding the Field Challenge | [Documentation](#)

Crop Classification from Satellite Time-Series Data

June 2025 – July 2025

- ➔ Engineered hybrid features by combining statistical and spectral vegetation indices with deep time-series embeddings from a PatchTST model.
- ➔ Upgraded an ensemble classification pipeline integrating XGBoost, Logistic Regression, and SVM for robust multi-crop prediction (cocoa, rubber, oil palm).
- ➔ Achieved Top 10% rank (10/100 teams) on Zindi leaderboard with a private log loss of 0.7473.
- ➔ Generated final outputs as predicted crop type labels and probabilities, delivered in CSV format for analysis and submission.

### AI Coder Pro: Visual Studio Code Extension | [Link](#)

July 2025

- ➔ Built AI-powered Visual Studio Code extension using TypeScript and Node.js with Together AI integration
- ➔ Implemented multi-agent tools for code review, bug detection, refactoring, and automated test generation
- ➔ Created chat interface with memory capabilities and file upload functionality for enhanced user experience
- ➔ Tools Used.: TypeScript, Node.js, Artificial Intelligence, Natural Language Processing.

## CERTIFICATIONS & ACHIEVEMENTS

- Won prize money of USD 500 for most efficient solution.
- Earned Top 10% rank (10/100 teams) in the Amini GeoFM Decoding the Field Challenge
- Active contributor to open-source Machine Learning projects on GitHub
- Excel Business Intelligence – [Link](#)      SQL Beginner to Advanced – [Link](#)      Python Beginner to Advanced – [Link](#)